

---

# ShorterBetter: Guiding Reasoning Models to Find Optimal Inference Length for Efficient Reasoning

---

Jingyang Yi   Jiazheng Wang   Sida Li

Data Science Institute, The University of Chicago  
Chicago, IL 60637

{jingyang22, jiazhengw, listar2000}@uchicago.edu

## Abstract

Recent models such as OpenAI o1 and DeepSeek-R1 have demonstrated strong performance on reasoning-intensive tasks by generating extended Chain-of-Thought (CoT) traces. While longer reasoning helps with thorough exploration of solution paths for complex problems, it also often leads to inefficient and redundant outputs—a phenomenon commonly described as *overthinking*. In this paper, we propose ShorterBetter, a simple yet effective reinforcement learning method that enables reasoning models to learn their own optimal CoT lengths without manual supervision. We define the *Sample Optimal Length* (SOL) as the length of the shortest correct response among multiple generations, which serves as a dynamic reward signal to guide the model toward efficient reasoning. Applied to DeepSeek-Distill-Qwen-1.5B/7B as base models, ShorterBetter achieves 50%-80% reduction in output lengths in both in-domain and out-of-domain reasoning tasks while maintaining accuracy. Our reasoning trace analysis shows that ShorterBetter refines the structure of the reasoning traces by reducing unnecessary repetition, excessive self-verification, and over-exploration of alternatives.<sup>1</sup>

## 1 Introduction

Improving the reasoning capabilities of Large Language Models (LLMs) has recently led to substantial performance gains in highly structured and complex tasks, such as mathematical problem-solving (Shao et al., 2024) and competitive programming (Jiang et al., 2024). Exemplified by OpenAI o1 (OpenAI, 2024) and DeepSeek-R1 (DeepSeek-AI, 2025), these reasoning models are incentivized to generate extended Chain-of-Thought (CoT) steps during inference. Such extended reasoning often includes detailed intermediate steps, reflective self-verifications, and exploration of alternative problem-solving strategies. These advancements are widely believed to enable LLMs to exhibit human-like *System-2 slow thinking* (Evans, 2003; Qu et al., 2025; Li et al., 2025), wherein significant test-time computational resources are deliberately allocated in exchange for improved accuracy and robustness on challenging reasoning benchmarks (Jain et al., 2024; Rein et al., 2024).

However, recent studies have shown that excessively long reasoning processes can be counter-productive, leading not only to computational inefficiencies but also to the so-called *overthinking phenomenon* (Chen et al., 2024; Luo et al., 2025a), where additional reasoning steps yield diminishing returns or even negatively impact model performance (Su et al., 2025).<sup>2</sup> For example, we identify several common problematic patterns in lengthy reasoning traces: (1) verbose repetitions without new insights, (2) early stopping of promising paths, (3) continued exploration of alternative solutions

---

<sup>1</sup>We will open source our training script, model weights and analysis pipelines.

<sup>2</sup>We provide an experiment in Appendix B.3 demonstrating that incorrect reasoning traces often have longer output lengths.




Figure 1: **ShorterBetter exhibits strong length reduction performance (1.5B models).** We report accuracy and output length of five sets of models on various tasks. Models include DeepSeek-R1-Distill-Qwen-1.5B, Qwen2.5-1.5B-Instruct (non-reasoning model), ShorterBetter-1.5B ( $\alpha=1$  and  $\alpha=2$ ) and Training Efficient-1.5B (Arora and Zanette, 2025). Our method achieves significant length reduction compared to all baseline reasoning models while maintaining accuracy.

despite having arrived at correct answers, and (4) collapse of reasoning coherence following very long outputs.<sup>3</sup> These observations indicate that lengthy reasoning traces are highly compressible, motivating techniques producing more concise reasoning chains that improve computational efficiency while preserving overall performance.

Existing approaches targeting efficient LLM reasoning primarily focus on *explicitly controlling reasoning length*, for instance, by imposing a *length budget* based on user specifications (L1; Aggarwal and Welleck (2025)), or by incentivizing shorter reasoning paths through comparisons with reference models (O1-Pruner; Luo et al. (2025a)). Further discussions on related approaches are provided in Section 2. In contrast to these explicit control strategies, our work is based on the hypothesis that reasoning models implicitly possess an *optimal reasoning length* (OL), which is dependent on model capacity and problem difficulty. OL thus represents an in-between "sweet spot": reasoning that is shorter might contain insufficient analysis, while reasoning that is longer risks redundancy and overthinking. This hypothesis has recently undergone theoretical formulation (Wu et al., 2025) and empirical support (Su et al., 2025). Since OL is intrinsic and not directly observable, we approximate it with the *Sample Optimal Length* (SOL): for each problem, we generate multiple outputs and define its SOL as the length of the shortest correct response. Leveraging SOL as a practical proxy for OL, we propose ShorterBetter, a reinforcement learning method that teaches the model to dynamically identify and steer towards efficient reasoning lengths during inference:

In summary, our contributions are as follows.

1. We propose ShorterBetter, a novel reinforcement learning method that guides reasoning model to autonomously learn its optimal reasoning lengths by defining and using Sample Optimal Length as a dynamic reward signal.
2. We empirically demonstrate that ShorterBetter, when applied to 1.5B/7B base models, achieves substantial (50%-80%) reductions in output token length on both in-domain mathematical reasoning tasks and diverse out-of-domain benchmarks, while maintaining accuracy.
3. We introduce novel reasoning trace analysis frameworks to demonstrate how ShorterBetter enhances reasoning efficiency. Specifically, the model reduces verbosity after finding the correct answer and refactors the overall reasoning traces to adopt more confident and targeted reasoning strategies.

<sup>3</sup>We include concrete examples for each of the aforementioned patterns in Appendix C.

## 2 Related Work

**Reasoning in LLMs** Since Wei et al. (2022) demonstrated that CoT prompting enhances LLM reasoning, following research has explored ways to scale test-time computation—such as parallel sampling (Wang et al., 2022) and iterative refinement (Madaan et al., 2023)—to improve performance. This led to Large Reasoning Models (LRMs) that autonomously generate extended CoTs through reinforcement learning-based post-training (Zhang et al., 2025). Examples include OpenAI’s o1 (OpenAI, 2024), DeepSeek-R1 (DeepSeek-AI, 2025), and Qwen-QwQ (Yang et al., 2024). These models internalize the ability to deliberate over multiple steps before producing final answers, achieving significant performance gains on reasoning-intensive benchmarks (Jimenez et al., 2023; Jain et al., 2024; Gao et al., 2024). Our work addresses a complementary *challenge*: how to dynamically optimize reasoning efficiency without sacrificing accuracy?

**The Overthinking Phenomenon** This challenge, which was initially observed as verbosity and increased computational costs in CoT prompting (Nayab et al., 2024; Han et al., 2024), has become more pronounced with the advent of LRMs—following the identification and understanding of the *overthinking phenomenon* (Chen et al., 2024; Su et al., 2025). This phenomenon is typically understood through two primary lenses. Firstly, overthinking is recognized as a significant source of computational inefficiency. In this view, LRMs expend excessive resources generating overly lengthy and often redundant reasoning traces, especially for simple problems (Lee et al., 2025). This inefficiency manifests as increased token usage, higher latency, and greater computational expenditure (Aytes et al., 2025), and might be exploited adversarially to induce slowdowns (Kumar et al., 2025).

Secondly, overthinking is increasingly understood as a factor that can directly impair practical performance. Fatemi et al. (2025) find that excessive reasoning can degrade logical correctness and introduce irrelevant information; Su et al. (2025) reveal that beyond a certain threshold, longer reasoning does not necessarily improve and can even degrade accuracy, as models may misjudge problem difficulty and introduce compounding errors; similarly, Wu et al. (2025) empirically and theoretically demonstrate a U-shaped performance curve with CoT length, showing that excessively long reasoning paths are susceptible to noise and error accumulation, ultimately reducing correctness. This performance degradation has also been observed in agentic tasks where prolonged internal deliberation correlates with lower task success rates (Cuadron et al., 2025).

**Efficient Reasoning with RL** Many recent and contemporaneous methods have been proposed to control the reasoning length in LRMs. Here, we specifically highlight RL-based post-training strategies directly relevant to our method, leaving readers to the survey by Qu et al. (2025) for an extensive overview. One category of methods uses an explicit *length budget*. For instance, L1 (Aggarwal and Welleck, 2025) treats the budget as a user-specified parameter, while DAST (Shen et al., 2025) adjusts it adaptively based on problem difficulty. These methods then use RL to reward generations that adhere closely to this budget. Other methods design length rewards differently. O1-Pruner (Luo et al., 2025a), for example, compares generated lengths against those from a reference model. More in line with our proposal of sampling multiple responses, Kimi 1.5 (Team, 2025) and Training Efficient (Arora and Zanette, 2025) utilize the shortest and average lengths, respectively, among *all* sampled responses as reward baselines. In contrast, ShorterBetter innovatively anchors its reward function to the length of the shortest *correct* response—a design inspired by recent studies on the existence of optimal reasoning lengths in LRMs (Wu et al., 2025; Lee et al., 2025; Su et al., 2025).

## 3 Method

In this work, we aim to guide the reasoning model to *think for the right amount of time*, which requires knowing its *optimal reasoning length* (OL)—a prompt-dependent quantity we formally define in Section 3.1. However, calculating OL in practice is challenging due to the absence of precise measures for both problem difficulty and model capability. To address this, we propose in Section 3.2 a sample-based approximation of OL, yielding a practical length reward. Finally, we integrate this reward design with Group Relative Policy Optimization (GRPO; Shao et al. (2024)) in Section 3.3, resulting in a simple yet effective reinforcement learning method that encourages the model to autonomously discover its optimal reasoning lengths through on-policy learning.




Figure 2: **Flowchart for the ShorterBetter pipeline.** We detail the reward design in Section 3.2 and the GRPO algorithm in Section 3.3. Reward values in display are rescaled for better illustration.

### 3.1 Optimal Reasoning Length (OL): Formal Definition

**Problem Setup** Let  $\mathcal{M}_\theta$  be our policy model (i.e. a base LRM) with parameters  $\theta$ , and  $\mathcal{D}$  be the training dataset. Each problem  $(x_i, y_i^*) \in \mathcal{D}$  consists of a prompt  $x_i$  and a reference response  $y_i^*$ . Given  $x_i$ , for each *rollout* the model generates a response  $y = (y^1, \dots, y^t)$  autoregressively from  $y \sim p_\theta(\cdot | x_i)$ , where  $y^k$  denotes its  $k$ -th token and  $\ell(y) = t$  is its length. For reasoning tasks,  $y$  typically includes long CoT tokens (often enclosed in `</think>` tags) followed by a final answer.

To evaluate response quality, we assume access to a verifier  $\mathcal{I}$  that assigns a normalized *correctness score*  $s_{\mathcal{I}}(y, y_i^*) \in [0, 1]$ , where higher scores indicate better alignment with  $y_i^*$ . For fixed correctness threshold  $c \geq 0$  and probability threshold  $\epsilon \geq 0$ , the optimal reasoning length is defined as

$$\ell_{c,\epsilon}^*(x_i; \theta) := \min\{\ell(y) : y \in \mathcal{Y}_{c,\epsilon}(x_i; \theta)\}, \quad \text{where}$$

$$\mathcal{Y}_{c,\epsilon}(x_i; \theta) := \{y : p_\theta(y | x_i) \geq \epsilon \wedge s_{\mathcal{I}}(y, y_i^*) \geq c\}.$$

Thus,  $\ell_{c,\epsilon}^*(x_i; \theta)$  represents the minimal length at which the model reliably (probability above  $\epsilon$ ) produces a sufficiently correct (score above  $c$ ) response. Intuitively, OL is longer for more difficult problems, as the feasible set  $\mathcal{Y}_{c,\epsilon}(x_i; \theta)$  shrinks due to fewer correct responses; conversely, OL becomes shorter when the model  $\mathcal{M}_\theta$  is capable of generating concise and correct responses more consistently. This intuition aligns with theoretical insights presented in Wu et al. (2025), despite differences in their specific formulation.

**Remark** The above formulation applies broadly to both closed-ended and open-ended reasoning tasks. For the remainder of this paper, we narrow our focus to tasks with *verifiable answers*, such as mathematical and coding problems. These tasks naturally come with a binary verifier that outputs

$$\mathbb{I}(y = y_i^*) := s_{\mathcal{I}}(y, y_i^*) \in \{0, 1\},$$

indicating perfect correctness (1) or incorrectness (0). Consequently, we adopt the trivial threshold choice  $c = 1$  and use the shorthand notation  $\mathbb{I}(y = y_i^*)$  to denote correctness.

### 3.2 Sample Optimal Length (SOL): A Practical Approximation

Directly computing the optimal reasoning length defined in the previous subsection is practically infeasible, as it requires complete knowledge of the model’s output distribution. To address this, we introduce the *Sample Optimal Length* (SOL)—a practical, sample-based approximation of OL. Given a prompt  $x_i$ , our method generates  $n$  rollouts of candidate responses  $G(x_i) = \{y_1, y_2, \dots, y_n\}$  from the policy  $p_\theta(\cdot | x_i)$ . We then define the SOL for the group  $G(x_i)$  as:

$$\ell^{\text{SOL}}(G(x_i)) = \begin{cases} \min_{y_j \in G(x_i): \mathbb{I}(y_j = y_i^*)=1} \ell(y_j), & \text{if at least one response is correct,} \\ \frac{1}{n} \sum_{j=1}^n \ell(y_j), & \text{otherwise.} \end{cases}$$

Given the SOL, we define the following reward function for each response  $y_j \in G(x_i)$ :

$$r(y_j) = \alpha \cdot \mathbb{I}(y_j = y_i^*) - \beta \cdot |\ell(y_j) - \ell^{SOL}(G(x_i))|, \quad (1)$$

where  $\alpha, \beta > 0$  are hyperparameters balancing correctness reward and length penalty, respectively.

Our SOL-based reward has several attractive properties:

- ✓ **No human intervention required:** Unlike explicit length-budget approaches, SOL is entirely data-driven and does not require manual tuning of problem-specific length thresholds.
- ✓ **Adaptive to difficulty and model capability:** SOL naturally becomes longer for harder problems and shorter for easier ones, as it dynamically adjusts based on the sampled responses; on the other hand, as the model improves during training, SOL automatically updates to reflect the shortest correct reasoning length achievable by the current policy.
- ✓ **Balanced handling of failure cases:** When no correct responses exist within a sampled group, we set SOL to the average response length as a neutral baseline. Our approach avoids biasing the model toward either overthinking or underthinking when there is no correct signal.

We defer the detailed discussion on specific hyperparameter choices  $(\alpha, \beta, n)$  to Section 4, and the ablation studies on our reward design to Section 5.3. To see the training dynamics of ShorterBetter, please refer to Appendix B.2.

### 3.3 Policy Optimization via GRPO

To optimize the model under the SOL-based reward function, we adopt the Group Relative Policy Optimization (GRPO) algorithm (Shao et al., 2024). GRPO operates on groups of sampled trajectories (in our case, a group  $G(x_i)$  has  $n$  samples) and assigns within-group *advantage estimates* based on normalizing the rewards  $\mathbf{r} = [r(y_1), \dots, r(y_n)]$  calculated in Eq. (1). Specifically, we have

$$\hat{A}_{j,k} = \frac{r_j - \text{mean}(\mathbf{r})}{\text{std}(\mathbf{r})}$$

as the advantage for the  $k$ -th step (token) of the  $j$ -th sample, where ‘mean’ and ‘std’ calculate sample reward average and standard deviation. Let  $\mathcal{M}_{\theta_{\text{ref}}}$  be a frozen reference policy (usually a copy of the current model before a few gradient updates), and  $p_{\theta}^k := p_{\theta}(y_j^k \mid y_j^{<k}, x_i)$  denote the conditional probability on the first  $k - 1$  tokens, the GRPO loss is

$$\mathcal{L}_{\text{GRPO}}(\theta) = -\frac{1}{n} \sum_{j=1}^n \frac{1}{t_j} \sum_{k=1}^{t_j} \left[ W(\hat{A}_{j,k}, p_{\theta}^k, p_{\theta_{\text{ref}}}^k, \varepsilon) - \gamma \cdot \tilde{\mathbb{D}}_{\text{KL}}[p_{\theta}^k \parallel p_{\theta_{\text{ref}}}^k] \right], \quad (2)$$

where  $W$  denotes a clipped policy improvement objective incorporating the advantage  $\hat{A}_{j,k}$  and likelihood ratio between current and reference policies;  $\tilde{\mathbb{D}}_{\text{KL}}$  is a KL divergence approximator (Schulman, 2020);  $\varepsilon$  and  $\gamma$  are hyperparameters for the clipping threshold and KL penalty, respectively. We provide the full expression and details in Appendix A.1.

Policy updates are then performed by minimizing  $\mathcal{L}_{\text{GRPO}}(\theta)$  via gradient descent, encouraging the model to increase the likelihood of high-reward samples while constraining deviation from the reference policy. We choose GRPO for its compatibility with group-based reward designs and its empirical success in training LRMs (DeepSeek-AI, 2025; Face, 2025). Finally, we include a flowchart illustration of our full method in Fig. 2.

## 4 Experimental Setup

**Datasets** We fine-tune our base model on DeepScaleR-preview (Luo et al., 2025b), a dataset which comprises 40K mathematical problems drawn from AIME, AMC, Omni-MATH (Gao et al., 2024), and Still dataset (Min et al., 2024). This mixed dataset covers tasks from high-school math problems to Olympiad-level challenges. For in-domain evaluation, we holdout subsets of the training math datasets, containing 1K problems in total. Additionally, we assess the out-of-domain generalization of our method on a set of reasoning-intensive benchmarks with a total of 6K problems, including MathQA (Amini et al., 2019) MMLU (Hendrycks et al., 2021), Big-Bench Hard (BBH; Suzgun et al. (2022)) and coding benchmarks including LiveCodeBench (Jain et al., 2024), MBPP (Austin et al., 2021) and HumanEval (Chen et al., 2021).



**Base Models** We use DeepSeek-R1-Distill-Qwen-1.5B and 7B (DeepSeek-AI, 2025) as our base reasoning models and train our models from them. We will refer to our trained model as ShorterBetter-1.5B and ShorterBetter-7B in the subsequent sections.

**Evaluations** For evaluation, we generate one response for each problem. For each benchmark, we consider 3 metrics: (1) **Accuracy**, defined as the number of correct responses divided by the total number of responses; (2) **Output Length**, measured as the average number of tokens per benchmark; and (3) **Accuracy-Efficiency (AE) Score**, a composite metric balancing length reduction against accuracy preservation (introduced by Luo et al. (2025a)). AE Score rewards models that shorten outputs without sacrificing correctness, with higher values indicating better efficiency. For the full formulation and hyperparameter choices, see Appendix A.2.

**Implementation Details** We used Ver1 (Sheng et al., 2024), an open-source RL training library, for post-training. We fine-tune our 1.5B base model with a batch size of  $N = 64$  and a rollout group size of  $n = 8$ . Training was conducted for 300 steps with a fixed learning rate of  $1 \times 10^{-6}$ . Since the 7B variant converges faster due to its better model capacity, we trained for only 200 steps. Due to computational constraints, we set a maximum inference length of 6K tokens for the 1.5B model and 5K tokens for the 7B model during training. (We set 16K maximum tokens budgets during evaluation) These constraints only affect the early stages of training, as the mean inference lengths quickly drop below the limits. To study the effect of hyperparameters, we compared two 1.5B variants with  $\alpha = 1$  and  $\alpha = 2$ , while keeping  $\beta = 0.001$  fixed. For the GRPO training configuration, we set  $\gamma = 0.001$  and  $\varepsilon = 0.2$  in Eq. (2).

**Baselines** To assess the efficiency of our method, we compare our trained models against 4 baselines:

- **DeepSeek-R1-Distill-Qwen-1.5B/7B** (DeepSeek-AI, 2025): Models obtained by distillation from the Qwen series using larger DeepSeek-R1 models. We refer to these models as Distill-1.5B and Distill-7B in the following sections.
- **Qwen2.5-1.5B/7B-Instruct** (Team, 2024): Instruction-tuned models from the Qwen 2.5 family of large language models. We include these non-reasoning models as a contrast, i.e., fast thinking without deliberate reasoning. We denote them as Instruct-1.5B and Instruct-7B hereafter.
- **Training Efficient** (Arora and Zanette, 2025): We evaluated the 1.5B/7B open-source models released in this work, which have the same base models as ours.
- **O1-Pruner** (Luo et al., 2025a): We evaluate the 7B open-source model proposed in this work,<sup>4</sup> which is fine-tuned from the Marco-o1 base model (Zhao et al., 2024).

## 5 Results and Discussion

In this section, we present and analyze the effectiveness of our proposed method. We compare the performance of our trained models against the baselines across a range of in-domain math tasks as well as out-of-domain benchmarks.<sup>5</sup> We then provide an in-depth analysis of how our method changed the structure of the reasoning traces.

### 5.1 Experimental Results

**ShorterBetter significantly reduces inference length while maintaining performance** In Fig. 1, we compare the performance of the ShorterBetter-1.5B against Distill-1.5B, Instruct-1.5B and Training Efficient-1.5B. We report results for two variants of ShorterBetter-1.5B, with hyperparameters  $\alpha = 1$  and  $\alpha = 2$ , respectively. In in-domain math tasks, averaging over all benchmarks, the two models achieve inference length reductions of 77.6% and 79.2%, with corresponding changes in accuracy of +2.5 % and +3.2 %, respectively. In out-of-domain tasks, averaging over all benchmarks, the two variants achieve average length reductions of

<sup>4</sup>Since this work does not provide a 1.5B model, we will only use 7B version as the baseline.

<sup>5</sup>In-domain and out-of-domain tasks are defined relative to ShorterBetter’s training dataset. See more discussions in Appendix A.3

	Distill		Instruct		ShorterBetter ( $\alpha = 2$ )			Training Efficient			O1-pruner		
	Acc.	Len.	Acc.	Len.	Acc.	Len.	AE Score $\uparrow$	Acc.	Len.	AE Score $\uparrow$	Acc.	Len.	AE Score $\uparrow$
<i>In-domain Tasks</i>													
AIME	36.7	11382	16.7	1172	53.3	5288	1.89	50.0	8821	1.31	6.7	4843	0.28
Olympiad	52.0	7762	38.2	944	50.7	3410	0.44	53.9	5544	0.40	29.6	3884	0.05
AMC	66.3	7414	51.8	1084	75.9	2580	1.09	72.3	4596	0.65	30.1	4304	0.31
Minerva	40.8	5034	38.2	1367	44.1	1341	0.98	40.4	2796	0.40	23.5	5277	0.55
<i>Out-of-domain Tasks</i>													
MathQA	83.3	3442	68.9	2585	85.4	980	0.79	78.8	1350	0.34	32.8	6321	0.14
BBH	68.3	2126	67.3	953	63.1	535	0.37	62.1	798	0.17	55.4	5431	-0.35
MMLU	66.3	1528	76.0	269	68.6	567	0.73	69.3	891	0.55	72.3	5979	0.31
LiveCodeBench	50.6	8683	24.7	189	49.5	5228	0.29	49.7	7589	0.04	21.6	9726	0.20

Table 1: **Performance comparison of 7B models across diverse tasks.** For each task, we highlight the best-performing results across all evaluated models. The AE Scores are computed using hyperparameters  $\varphi = 1$ ,  $\eta = 3$ , and  $\theta = 5$ . The proposed ShorterBetter model consistently achieves the highest AE Scores across both in-domain and out-of-domain benchmarks. For experimental details and additional results on coding tasks, refer to Appendix A and Appendix B.1.

77.0 % and 75.5 % , with corresponding changes in accuracy of  $-0.8$  % and  $-1.6$  % , respectively.<sup>6</sup> These results also indicate that our method achieves similar performance in terms of length reduction and accuracy preservation with different values of correctness reward hyperparameter  $\alpha$ .

In Table 1, we provide a comprehensive comparison between ShorterBetter-7B ( $\alpha = 2$ ) and four 7B baseline models.<sup>7</sup> Consistent with our 1.5B model evaluation, we assess all 7B models on both in-domain and out-of-domain tasks. For Distill-7B and Instruct-7B, we only report accuracy and output length as primary evaluation metrics since they didn’t undergo length reduction training. For models trained with length-reduction techniques, we additionally report the AE Score, which quantifies performance improvements relative to their respective base models. (See a formal definition of AE Score in Appendix A.2)

Our 7B model achieves substantial reductions in output length, up to 70%, while maintaining competitive accuracy compared to the Distill-7B baseline. In in-domain math tasks, averaging over all benchmarks, ShorterBetter-7B attains a length reduction of 62.1% alongside an accuracy improvement of +7.1%. In out-of-domain tasks, averaging over all benchmarks, the model achieves a length reduction of 62.3% with an accuracy change of -0.48%. Moreover, ShorterBetter-7B attains the highest AE Scores in all evaluated tasks, outperforming Training Efficient and O1-Pruner.

Compared to prior length-reduction methods, ShorterBetter achieves stronger length-reduction capability. We attribute this performance to the design of SOL, which gives clear signals to guide the model to learn concise and robust reasoning traces during reinforcement learning. For more discussions on the effectiveness of SOL, please refer to Appendix B.2. In the next section, we show that our method refines the structure of the reasoning traces relative to the base model.

## 5.2 Reasoning Structure Analysis

In this section, we present two novel frameworks for analyzing reasoning structure, which illustrate how ShorterBetter compresses its reasoning traces and complement the accuracy and token-length metrics introduced earlier. We employ Gemini 2.5 Flash Preview (Google, 2025) as an automated judge (details in Appendix D.1) to breakdown and compare the reasoning traces of ShorterBetter-7B and Distill-7B, shedding lights on how ShorterBetter optimizes for reasoning efficiency.

**Output Length after First Appearance of Correct Answer** We use LLM-as-a-judge to identify the first occurrence of the correct answer in the reasoning trace for questions where both

<sup>6</sup>Due to space constraint, we only present LiveCodeBench here and place other coding benchmarks in Fig. 5 and Table 2 in the appendix

<sup>7</sup>For a fair comparison, we compute the AE Scores of O1-Pruner using its corresponding base model, Marco-o1-7B, whereas the base models for ShorterBetter and Training Efficient are both Distill-7B. Due to space limitations, the performance of Marco-o1-7B is placed in Appendix B.4




Figure 3: **Reasoning trace structure analysis.** **Left:** Percentage of remaining output after first appearance of the correct answer; **Right:** Token-level distribution of reasoning categories on the MATH benchmark.

ShorterBetter-7B and Distill-7B provide correct answers. We conduct this analysis across four benchmarks: AMC, MATH, Olympiad, and Minerva. For each model, we calculate what percentage of its total output (in tokens) appears after the first time that the correct answer is stated. The results are shown in Fig. 3 (Left), where higher percentage represents more reasoning outputs after the model finds the correct answer. We observe that the percentages of tokens appearing after the first correct answer decreases by 15%, 20%, 17%, and 19% for AMC, MATH, Olympiad, and Minerva benchmarks, respectively. This suggests that ShorterBetter-7B reaches conclusions more efficiently, using fewer tokens on extra reasoning once it reaches the correct answer, and we will explore more detailed reasoning structure refinement in the next section.

**Breakdown of Reasoning Traces** To better understand how ShorterBetter improves reasoning efficiency, we conduct a fine-grained analysis of the reasoning trace structure. Specifically, we develop an LLM-as-a-judge framework to assign functional categories to each sentence in the generated traces from ShorterBetter-7B and Distill-7B, and compute the token-level proportion of each reasoning category. We begin by carefully inspecting common patterns in Chain-of-Thought reasoning and curating a set of mutually exclusive categories, each assigned at the sentence level:

- **Pivotal Reasoning:** Steps that directly correspond to a specific part of the final solution.
- **Productive Elaboration & Calculation:** Necessary calculations, deductions, or planning that support a pivotal step but aren’t restated in the final summary.
- **Exploring Alternatives:** Attempts to try different approaches or check other methods, even if they aren’t used in the final solution.
- **Verification & Self-Correction:** Checks and corrections of earlier results to catch errors.
- **Non-Substantive Statement:** Redundant comments, filler, or trivial rephrasing that doesn’t advance the solution.

We then use LLM-as-a-judge to label each sentence with the most appropriate category.<sup>8</sup> Finally, we compute the proportion of each category by counting tokens using the tested model’s native tokenizer, enabling a consistent token-level comparison across models. In Fig. 3 (Right), we show the distribution of reasoning categories for both models on the MATH benchmark. Compared to Distill-7B, ShorterBetter-7B exhibits notable structural shifts. Specifically, the proportions of *Pivotal Reasoning* and *Productive Elaboration & Calculation* increase by 15.6% and 6.8%, respectively—indicating a denser concentration of meaningful, solution-oriented reasoning. Simultaneously, the share of *Non-Substantive Statements* decreases by 4.8%, suggesting that ShorterBetter-7B produces fewer filler or redundant tokens. We also observe reductions in *Verification & Self-Correction* (-14.1%) and *Exploring Alternatives* (-3.9%), implying that the model is more decisive and less prone to unnecessary detours. Altogether, these shifts support the conclusion that ShorterBetter promotes more confident and streamlined reasoning trajectories. This finding aligns with our analysis in the previous section and may help address the early stopping of promising paths issue mentioned in Section 1.

<sup>8</sup>For transparency and reproducibility, our prompt and examples of each category are provided in Appendix D.2






Figure 4: **Length and accuracy trends of the ablation experiments.** **Left:** Discard the entire design of SOL and partially discard SOL. Denoted as *Discard* and *Partial* in the plot. **Right:** Remove the correct rewarding term in the reward function but keep the entire design of SOL, i.e., set  $\alpha = 0$ . Both ablation experiments are conducted on Distill-7B.

While our reasoning structure analysis in Section 5.2 provides valuable insights, we encourage future works to develop more rigorous quantitative metrics for evaluating reasoning structure, which could further validate and extend our findings.

### 5.3 Ablation Studies

**Naively adopting the shortest response as target length leads to training collapses** We conducted two ablation experiments: (1) Discarding the entire design of ShorterBetter: remove the correctness term in the reward function and always choose the length of the shortest response as the reward signal. (2) Partially change the design of SOL : when there is no correct response, select the length of the shortest response as the SOL, instead of using the average length as in Eq. (1).<sup>9</sup> Note that we still keep the correctness reward term in experiment (2). In Fig. 4 (left), model trainings under both ablation experiment settings quickly collapse within 100 steps, suggesting that naively choosing the shortest response length as a reward can easily destabilize the training.

**Sample optimal length inherently combines correctness and conciseness** We conducted another experiment that keeps the entire design of SOL while discarding the correctness reward term in the reward function, i.e., set  $\alpha = 0$ . As shown in Fig. 4 (right), although there are drastic fluctuations in accuracy due to the absence of explicit correct rewards, the training still runs effectively. This result suggests that the design of SOL inherently combines correctness and conciseness. Nonetheless, incorporating a correctness reward term remains essential to prevent degradation in model accuracy.

## 6 Conclusion and Future Directions

This paper presents ShorterBetter, a reinforcement learning method that enables reasoning models to dynamically optimize their CoT lengths without manual supervision. By introducing Sample Optimal Length as a self-supervised reward signal, our method achieves significant reductions in output length while maintaining accuracy across in-domain and out-of-domain benchmarks. We also present innovative analytical framework to show ShorterBetter’s reasoning structure improvement.

While ShorterBetter is currently showing strong empirical results on tasks with *verifiable answers*, we note that our general formulation in Section 3.1 naturally extends to open-ended settings with non-binary correctness scores. Another future direction is to apply and evaluate our methods on larger-scale reasoning models (Team, 2024). Finally, our reasoning trace analysis shows that different categories of reasoning behavior are compressed to different extents, which may inform the design of more fine-grained and behavior-aware reasoning optimization strategies.

<sup>9</sup>Note that in experiment (1), we replace SOL with sample minimal length and discard the correctness term in the reward function. The reward function is rewritten as  $r(y_j) = -\beta \left| \min_{y_i \in G(x_i)} \ell(y_i) - \ell(y_j) \right|$ . While in experiment (2), we still keep the correctness term in the reward function and set  $\alpha = 2$ .

## References

- Pranjal Aggarwal and Sean Welleck. L1: Controlling how long a reasoning model thinks with reinforcement learning. *arXiv preprint arXiv:2503.04697*, 2025.
- Aida Amini, Saadia Gabriel, Shanchuan Lin, Rik Koncel-Kedziorski, Yejin Choi, and Hannaneh Hajishirzi. MathQA: Towards interpretable math word problem solving with operation-based formalisms. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 2357–2367, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1245. URL <https://aclanthology.org/N19-1245>.
- Daman Arora and Andrea Zanette. Training language models to reason efficiently. *arXiv preprint arXiv:2502.04463*, 2025.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Simon A Aytes, Jinheon Baek, and Sung Ju Hwang. Sketch-of-thought: Efficient llm reasoning with adaptive cognitive-inspired sketching. *arXiv preprint arXiv:2503.05179*, 2025.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde De Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Xingyu Chen, Jiahao Xu, Tian Liang, Zhiwei He, Jianhui Pang, Dian Yu, Linfeng Song, Qiuzhi Liu, Mengfei Zhou, Zhuosheng Zhang, et al. Do not think that much for  $2+3=?$  on the overthinking of o1-like llms. *arXiv preprint arXiv:2412.21187*, 2024.
- Alejandro Cuadron, Dacheng Li, Wenjie Ma, Xingyao Wang, Yichuan Wang, Siyuan Zhuang, Shu Liu, Luis Gaspar Schroeder, Tian Xia, Huanzhi Mao, et al. The danger of overthinking: Examining the reasoning-action dilemma in agentic tasks. *arXiv preprint arXiv:2502.08235*, 2025.
- DeepSeek-AI. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning, 2025. URL <https://arxiv.org/abs/2501.12948>.
- Jonathan St BT Evans. In two minds: dual-process accounts of reasoning. *Trends in cognitive sciences*, 7(10):454–459, 2003.
- Hugging Face. Open r1: A fully open reproduction of deepseek-r1, January 2025. URL <https://github.com/huggingface/open-r1>.
- Mehdi Fatemi, Banafsheh Rafiee, Mingjie Tang, and Kartik Talamadupula. Concise reasoning via reinforcement learning. *arXiv preprint arXiv:2504.05185*, 2025.
- Bofei Gao, Feifan Song, Zhe Yang, Zefan Cai, Yibo Miao, Qingxiu Dong, Lei Li, Chenghao Ma, Liang Chen, Runxin Xu, et al. Omni-math: A universal olympiad level mathematic benchmark for large language models. *arXiv preprint arXiv:2410.07985*, 2024.
- Google. Gemini 2.5 flash. <https://blog.google/products/gemini/gemini-2-5-flash-preview/>, April 2025. URL <https://blog.google/products/gemini/gemini-2-5-flash-preview/>. Accessed: 2025-05-15.
- Tingxu Han, Zhenting Wang, Chunrong Fang, Shiyu Zhao, Shiqing Ma, and Zhenyu Chen. Token-budget-aware llm reasoning. *arXiv preprint arXiv:2412.18547*, 2024.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021.
- Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. Livecodebench: Holistic and contamination free evaluation of large language models for code. *arXiv preprint arXiv:2403.07*, 2024.

- Juyong Jiang, Fan Wang, Jiasi Shen, Sungju Kim, and Sunghun Kim. A survey on large language models for code generation. *arXiv preprint arXiv:2406.00515*, 2024.
- Carlos E Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. Swe-bench: Can language models resolve real-world github issues? *arXiv preprint arXiv:2310.06770*, 2023.
- Abhinav Kumar, Jaechul Roh, Ali Naseh, Marzena Karpinska, Mohit Iyyer, Amir Houmansadr, and Eugene Bagdasarian. Overthink: Slowdown attacks on reasoning llms. *arXiv e-prints*, pages arXiv-2502, 2025.
- Ayeong Lee, Ethan Che, and Tianyi Peng. How well do llms compress their own chain-of-thought? a token complexity approach. *arXiv preprint arXiv:2503.01141*, 2025.
- Zhong-Zhi Li, Duzhen Zhang, Ming-Liang Zhang, Jiaxin Zhang, Zengyan Liu, Yuxuan Yao, Haotian Xu, Junhao Zheng, Pei-Jie Wang, Xiuyi Chen, et al. From system 1 to system 2: A survey of reasoning large language models. *arXiv preprint arXiv:2502.17419*, 2025.
- Haotian Luo, Li Shen, Haiying He, Yibo Wang, Shiwei Liu, Wei Li, Naiqiang Tan, Xiaochun Cao, and Dacheng Tao. O1-pruner: Length-harmonizing fine-tuning for o1-like reasoning pruning. *arXiv preprint arXiv:2501.12570*, 2025a.
- Michael Luo, Sijun Tan, Justin Wong, Xiaoxiang Shi, William Tang, Manan Roongta, Colin Cai, Jeffrey Luo, Tianjun Zhang, Erran Li, Raluca Ada Popa, and Ion Stoica. Deepscaler: Surpassing o1-preview with a 1.5b model by scaling rl, 2025b. URL <https://pretty-radio-b75.notion.site/DeepScaleR-Surpassing-O1-Preview-with-a-1-5B-Model-by-Scaling-RL-19681902c1468005bed8ca303013a4e2>. Notion Blog.
- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement with self-feedback. *Advances in Neural Information Processing Systems*, 36:46534–46594, 2023.
- Yingqian Min, Zhipeng Chen, Jinhao Jiang, Jie Chen, Jia Deng, Yiwon Hu, Yiru Tang, Jiapeng Wang, Xiaoxue Cheng, Huatong Song, et al. Imitate, explore, and self-improve: A reproduction report on slow-thinking reasoning systems. *arXiv preprint arXiv:2412.09413*, 2024.
- Sania Nayab, Giulio Rossolini, Marco Simoni, Andrea Saracino, Giorgio Buttazzo, Nicolamaria Manes, and Fabrizio Giacomelli. Concise thoughts: Impact of output length on llm reasoning and cost. *arXiv preprint arXiv:2407.19825*, 2024.
- OpenAI. Learning to reason with language models, 2024. URL <https://openai.com/index/learning-to-reason-with-llms/>. Accessed: 2025-04-21.
- Xiaoye Qu, Yafu Li, Zhaochen Su, Weigao Sun, Jianhao Yan, Dongrui Liu, Ganqu Cui, Daizong Liu, Shuxian Liang, Junxian He, et al. A survey of efficient reasoning for large reasoning models: Language, multimodality, and beyond. *arXiv preprint arXiv:2503.21614*, 2025.
- David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R Bowman. Gpqa: A graduate-level google-proof q&a benchmark. In *First Conference on Language Modeling*, 2024.
- John Schulman. Approximating KL divergence, 2020. URL <http://joschu.net/blog/kl-approx.html>. Blog post.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Y Wu, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Yi Shen, Jian Zhang, Jieyun Huang, Shuming Shi, Wenjing Zhang, Jiangze Yan, Ning Wang, Kai Wang, and Shiguo Lian. Dast: Difficulty-adaptive slow-thinking for large reasoning models. *arXiv preprint arXiv:2503.04472*, 2025.

- Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. Hybridflow: A flexible and efficient rlhf framework. *arXiv preprint arXiv: 2409.19256*, 2024.
- Jinyan Su, Jennifer Healey, Preslav Nakov, and Claire Cardie. Between underthinking and overthinking: An empirical study of reasoning length and correctness in llms. *arXiv preprint arXiv:2505.00127*, 2025.
- Mirac Suzgun, Nathan Scales, Nathanael Schärli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung, Aakanksha Chowdhery, Quoc V Le, Ed H Chi, Denny Zhou, , and Jason Wei. Challenging big-bench tasks and whether chain-of-thought can solve them. *arXiv preprint arXiv:2210.09261*, 2022.
- Kimi Team. Kimi k1. 5: Scaling reinforcement learning with llms. *arXiv preprint arXiv:2501.12599*, 2025.
- Qwen Team. Qwen2.5: A party of foundation models, September 2024. URL <https://qwenlm.github.io/blog/qwen2.5/>.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171*, 2022.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.
- Yuyang Wu, Yifei Wang, Tianqi Du, Stefanie Jegelka, and Yisen Wang. When more is less: Understanding chain-of-thought length in llms. *arXiv preprint arXiv:2502.07266*, 2025.
- An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2. 5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.
- Chong Zhang, Yue Deng, Xiang Lin, Bin Wang, Dianwen Ng, Hai Ye, Xingxuan Li, Yao Xiao, Zhanfeng Mo, Qi Zhang, et al. 100 days after deepseek-r1: A survey on replication studies and more directions for reasoning language models. *arXiv preprint arXiv:2505.00551*, 2025.
- Yu Zhao, Huifeng Yin, Bo Zeng, Hao Wang, Tianqi Shi, Chenyang Lyu, Longyue Wang, Weihua Luo, and Kaifu Zhang. Marco-o1: Towards open reasoning models for open-ended solutions. *arXiv preprint arXiv:2411.14405*, 2024.

## ♣ Appendix: Table of Contents

<b>A. Technical and Implementation Details</b> .....	Page 13
<b>B. Additional Experiment Results</b> .....	Page 14
<b>C. Examples of Unwanted Reasoning Behaviors</b> .....	Page 17
<b>D. Analyzing Reasoning Traces with LLM-as-a-judge</b> .....	Page 20

### Broader Impact

This work contributes to the growing body of research on efficient reasoning in LLMs. By reducing redundant reasoning steps while maintaining accuracy, our method could help lower computational costs and improve the practicality of deploying LLMs in resource-constrained settings. Like many advancements in ML, this research may have indirect societal implications—for instance, more efficient models could reduce energy consumption or enable faster inference in applications such as education or automated decision-making. However, our focus is narrowly technical, and we do not make strong claims about specific real-world benefits or risks. As with any optimization technique, the broader impact will depend on how it is integrated into larger systems and deployed in practice.

## A Technical and Experimental Details

### A.1 Detailed Specification of GRPO Algorithm

Here we provide the complete expressions and details of the GRPO (DeepSeek-AI, 2025) objective (loss function) mentioned in Section 3.3:

$$\mathcal{L}_{\text{GRPO}}(\theta) = -\frac{1}{n} \sum_{j=1}^n \frac{1}{t_j} \sum_{k=1}^{t_j} \left[ W(\hat{A}_{j,k}, p_{\theta}^k, p_{\theta_{\text{ref}}}^k, \varepsilon) - \gamma \cdot \tilde{\mathbb{D}}_{\text{KL}}[p_{\theta}^k \| p_{\theta_{\text{ref}}}^k] \right].$$

In particular,  $W$  is a clipped policy improvement objective with full expression

$$W(\hat{A}_{j,k}, p_{\theta}^k, p_{\theta_{\text{ref}}}^k, \varepsilon) = \min \left( \frac{p_{\theta}^k}{p_{\theta_{\text{ref}}}^k} \hat{A}_{j,k}, \text{clip} \left( \frac{p_{\theta}^k}{p_{\theta_{\text{ref}}}^k}; 1 - \varepsilon, 1 + \varepsilon \right) \hat{A}_{j,k} \right),$$

where the clipping function  $\text{clip}(x; a, b) := \min(b, \max(a, x))$  prevents the updated parameters from deviating too much from the reference policy. For KL divergence, Schulman (2020) proposes the following approximator

$$\tilde{\mathbb{D}}_{\text{KL}}[p_{\theta}^k \| p_{\theta_{\text{ref}}}^k] := \frac{p_{\theta_{\text{ref}}}^k}{p_{\theta}^k} - \log \frac{p_{\theta_{\text{ref}}}^k}{p_{\theta}^k} - 1.$$

Finally, while the above formulation considers multiple rollouts for a single prompt  $x_i$ , in practice we compute the GRPO loss and update the policy parameters over a batch of  $N$  prompts at each step.

### A.2 Detailed Formulation for Accuracy-Efficiency (AE) Score

The **Accuracy-Efficiency (AE) Score**, introduced by Luo et al. (2025a), provides a composite metric to evaluate whether a model improves inference efficiency—i.e., reduces output length—without sacrificing accuracy. It is defined as follows:

$$\text{AE Score} = \begin{cases} \varphi \cdot \Delta \text{Length} + \eta \cdot |\Delta \text{Acc}|, & \text{if } \Delta \text{Acc} \geq 0 \\ \varphi \cdot \Delta \text{Length} - \theta \cdot |\Delta \text{Acc}|, & \text{if } \Delta \text{Acc} < 0 \end{cases}$$

Here,  $\Delta \text{Length}$  and  $\Delta \text{Acc}$  denote the percentage changes in output length and accuracy, respectively, of the evaluated model relative to its base model. Positive AE Scores indicate that the model generates shorter outputs while maintaining or improving accuracy, while negative scores penalize accuracy degradation. In our experiments, we use the same hyperparameters as in Luo et al. (2025a):



- $\varphi = 1$  (weight on length reduction)
- $\eta = 3$  (bonus for accuracy gains)
- $\theta = 5$  (penalty for accuracy drops)

The asymmetric design (with  $\theta > \eta$ ) penalizes accuracy drops more heavily than gains, reflecting the practical preference for avoiding performance degradation.

### A.3 Experimental Details

For evaluation, we categorize tasks as either in-domain or out-of-domain, based on the training data used for ShorterBetter. Our method is trained on the DeepScaleR-Preview dataset (Luo et al., 2025b), a composite benchmark that includes AIME, AMC, Olympiad-Bench, MATH, Minerva, etc. Accordingly, held-out subsets from these datasets are considered in-domain. In contrast, tasks such as MathQA, BBH, MMLU, LiveCodeBench, MBPP, and HumanEval—which are not included in the training corpus—are treated as out-of-domain.

Other length-reduction baselines are trained on different datasets. Specifically, Training Efficient (Arora and Zanette, 2025) is trained on MATH, CN-K12, AIME, AoPS, and Olympiad-Bench, while O1-Pruner (Luo et al., 2025a) is trained solely on MATH. There exists partial overlap in training data between these baselines and ShorterBetter. Nevertheless, the out-of-domain benchmarks defined in this work are held out from all models’ training data and thus serve as a fair test of generalization.

For model comparisons in Fig. 1 and Table 1, we directly downloaded the open-source models of Training Efficient 1.5B/7B and O1-Pruner 7B from HuggingFace for evaluation. For O1-Pruner, since the paper does not provide a 1.5B version, we didn’t include it in Fig. 1. For Training Efficient, we used the 1.5B and 7B versions both with hyperparameters  $\alpha = 0.4$  (following the original paper). During all evaluations, the temperature was set to 0.9. All problems are answered once, and the accuracy is calculated as the fraction of correct answers among all answers.

### A.4 Computational Resources

We trained our ShorterBetter-1.5B with 4 A100 GPUs, using approximately 16 hours. We trained ShorterBetter-7B with 8 A100 GPUs, using less than 12 hours.

## B Additional Experiment Results

### B.1 Coding Benchmarks

We report the result of the coding experiment in Table 2 and Fig. 5.

Task	1.5B Models				7B Models				
	R1-Distill	Instruct	ShorterBetter	Training Efficient	R1-Distill	Instruct	ShorterBetter	Training Efficient	O1-pruner
MBPP	<b>0.544</b> (4151)	0.409 (124)	0.476 (1146)	0.506 (2402)	<b>0.659</b> (2928)	0.609 (150)	0.632 (1055)	0.607 (2113)	0.494 (495)
HumanEval	<b>0.720</b> (5123)	0.463 (263)	0.628 (1551)	0.665 (3280)	0.854 (3841)	0.811 (398)	0.799 (1318)	<b>0.866</b> (2601)	0.829 (254)
LiveCodeBench	<b>0.260</b> (10204)	0.089 (262)	<b>0.260</b> (4861)	0.244 (7792)	<b>0.506</b> (8683)	0.247 (189)	0.495 (5228)	0.216 (9726)	0.483 (1847)

Table 2: Coding performance across model sizes and training methods. Each entry reports the accuracy rate and the average inference length in tokens (shown in parentheses). Both ShorterBetter models were fine-tuned with  $\alpha = 2$ . Baseline models are from Arora and Zanette (2025). Bold numbers indicate best accuracy for each model size category.

### B.2 Training dynamics of ShorterBetter

In Fig. 6 we show the output length (in tokens) trends and accuracy trends of ShorterBetter-1.5B and ShorterBetter-7B during GRPO training. For clearer illustration, we present the first 200 steps of the two models.




Figure 5: Model performance on coding benchmarks.




Figure 6: Training dynamics of ShorterBetter-1.5B and ShorterBetter-1.5B models.

In Fig. 7, we illustrate how the model learns to identify optimal inference lengths throughout the training. For each sample generation ( $n = 8$ ) of a given problem, we compute the SOL and take the batch-wise average. We also track the mean deviation of individual output from their respective SOL, defined as  $d_i = \frac{\ell(y_i) - \ell^{SOL}(G(x_i))}{\ell^{SOL}(G(x_i))}$ . For instance, a deviation of 0.2 at a given training step indicates that the generated outputs deviate from the optimal length by 20% on average. The plots show that both the **SOL** and the mean deviation steadily decrease over training, particularly within the first 100 steps. This trend suggests that reinforcement learning is successfully guiding the model to produce outputs that cluster towards the optimal length.

### B.3 Incorrect reasoning traces are significantly longer

In Fig. 8, we conducted an experiment using 150 problems (30 each from AIME, AMC, MATH, Olympiad-bench, and Minerva). For each problem, we generated 10 responses with the Distill-1.5B/7B models, then filtered out problems where all responses were correct or all responses were wrong. We then computed the difference between the average lengths of correct and incorrect responses to form a length gap distribution. The results show that, on average, for the same problem,




Figure 7: SOL trends and mean deviations during training

incorrect responses are significantly longer than correct ones, by 5,541 tokens on DeepSeek-R1-Distill-Qwen-1.5B and 5,910 tokens on DeepSeek-R1-Distill-Qwen-7B. This suggests a clear correlation between reasoning errors and unnecessarily long CoT traces in Large Reasoning Models.

#### Length Differences Between Correct and Wrong Outputs

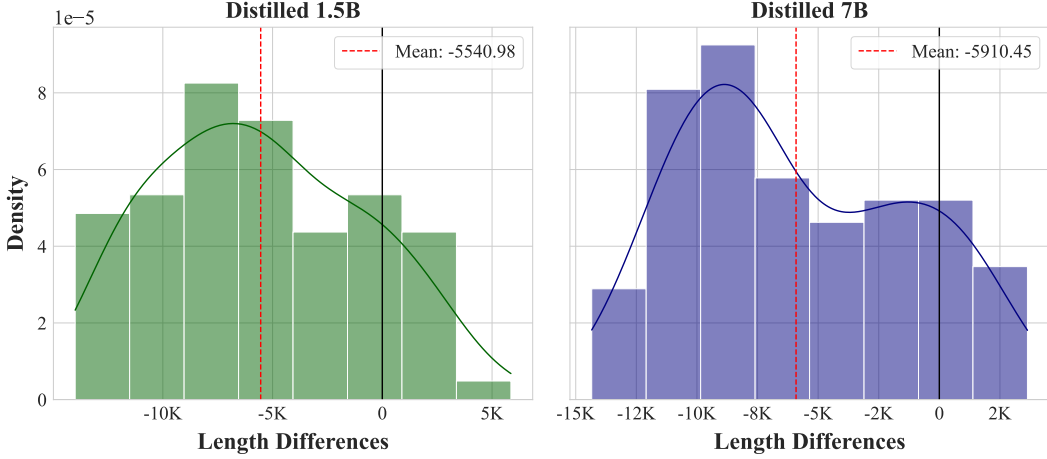


Figure 8: Length gap distributions for the Distill-1.5B and 7B models.

#### B.4 Performance of Marco-o1

Here we report the performance of Marco-o1 on various benchmarks, which is used to calculate the AE Scores of O1-Pruner in Table 1.

Dataset	Accuracy	Length
AIME	0.067	6711
Olympiad	0.317	6325
AMC	0.301	6210
Minerva	0.217	7583
MathQA	0.331	7730
BBH	0.623	6829
MMLU	0.701	7650
LiveCodeBench	0.225	12856
HumanEval	0.841	317
MBPP	0.454	877

Table 3: Marco-o1 performance across datasets

## C Examples of Unwanted Reasoning Behaviors

### C.1 Verbose Repetition and Self-Verification in Intermediate Steps

#### Prompt

Denali and Nate work for a dog walking business and are paid for each dog they walk. Denali is responsible for 16 dogs and Nate is responsible for 12 dogs. Under the company's new policy, they will be assigned or unassigned new dogs in groups of  $x$  dogs. The ratio of Denali's pay to Nate's pay would be the same if Denali started walking  $4x$  more dogs and Nate stayed at 12 dogs or if  $x$  of Nate's dogs were reassigned to Denali. Find  $x$  if  $x \neq 0$ .

#### Response (DeepSeek-R1-Distill-Qwen-1.5B)

...

**Wait**, perhaps the problem is not about the same pay ratio, but the same total number of dogs? But the problem says "the ratio of Denali's pay to Nate's pay would be the same." So it's about the pay ratio, not the number of dogs.

**Wait, another thought**: Maybe the ratio is in terms of the number of dogs walked, so if the ratio of their pay is same, and pay is per dog, then perhaps the number of dogs walked should be proportional.

**Wait, that is**, if Denali walks 16 dogs and Nate walks 12, the ratio of their pay is  $(16)/(12)$  if they are paid per dog.

...

**Wait**, so regardless of how I set up the ratio, I end up with  $x=0$ .

**But** since the problem states  $x \neq 0$ , perhaps my entire approach is wrong.

**Wait**, let me think differently.

**Perhaps** the ratio of their pay is the same when they are assigned  $x$  dogs, regardless of direction.

**Wait**, when they are assigned  $x$  dogs in groups of  $x$ , meaning  $x$  dogs are reassigned. So the ratio is same.

**Perhaps** it's the pay per dog ratio that's the same.

**Wait**, maybe Denali's pay per dog is same as Nate's.

**But** that would be same as each getting paid per dog.

**Wait**, but if they are reassigned, their pay per dog would be same.

**Wait**, perhaps we can think in terms of average.

**Wait**, I'm getting confused. Maybe I should go back to the original problem.

**\*\*Problem Statement:\*\*** ...

**But** we saw that this leads to  $x=0$ .

**Alternatively**, maybe the ratio is the other way around.

**Wait**, perhaps the ratio is the pay per dog ratio. So if the pay per dog is same, then the ratio is 1:1.

**But** in that case, both would have to walk same number of dogs. **But** the problem says that if  $x$  dogs are reassigned, it results in same ratio.

**But** in that case, if the pay per dog is same, then it's just the same number of dogs walked, which would imply that  $16 - x = 16 + 4x \Rightarrow -5x = 0 \Rightarrow x = 0$ , which is again invalid.

**Wait**, perhaps I need to approach this problem by assuming that the pay ratio remains the same.

...

## C.2 Early Stopping of Promising Reasoning Paths

### Prompt

The function  $f$  satisfies the functional equation

$$f(x) + f(y) = f(x + y) - xy - 1$$

for all real numbers  $x$  and  $y$ . If  $f(1) = 1$ , then find all integers  $n$  such that  $f(n) = n$ . Enter all such integers, separated by commas.

### Response (DeepSeek-R1-Distill-Qwen-1.5B)

...  
Now, let's find  $f(2)$ :  
Using the recursive formula, set  $x = 1$ :  
 $f(2) = f(1) + 1 + 2 = 1 + 1 + 2 = 4$   
Wait, is that right? Let me check.  
Wait, actually, according to the recursive formula  $f(x + 1) = f(x) + x + 2$ , so for  $x = 1$ : ...

We observe that the model starts to check for  $x = 1$  immediately after it decides to check  $f(2)$ , interrupting its own verification step.

## C.3 Continued Exploration of Alternative Solutions After Reaching Correct Answer

### Prompt

Find all values of  $x$  that satisfy the equation  $x = \sqrt{11 - 2x} + 4$ .

### Response (DeepSeek-R1-Distill-Qwen-1.5B)

...  
So, the solutions are  $x = 1$  and  $x = 5$ . But wait, I need to check if these solutions actually satisfy the original equation because sometimes squaring both sides can introduce extraneous solutions.  
First, let me check  $x = 1$ :  
Left side:  $1 = \sqrt{11 - 2(1)} + 4$  Calculate inside the square root:  $11 - 2 = 9$  So, square root of 9 is 3. Then  $3 + 4 = 7$  But left side is 1, so  $1 \neq 7$ . That means  $x = 1$  is not a valid solution.  
Now, check  $x = 5$ :  
Left side:  $5 = \sqrt{11 - 2(5)} + 4$  Calculate inside the square root:  $11 - 10 = 1$  Square root of 1 is 1. Then  $1 + 4 = 5$  So,  $5 = 5$ , which is true. Therefore,  $x = 5$  is a valid solution.  
But wait, earlier I thought that  $x$  must be greater than or equal to 4. Since  $x = 5$  is greater than 4, that's fine. But  $x = 1$  is less than 4, ...  
Wait, hold on. Let me double-check my earlier factorization and solutions.  
I had  $x^2 - 6x + 5 = 0$ , which factored into  $(x - 1)(x - 5) = 0$ , giving  $x = 1$  and  $x = 5$ .  
Plugging  $x = 5$  into the original equation worked, but plugging  $x = 1$  didn't. So,  $x = 1$  is extraneous.  
But is there any other step I should verify? Sometimes, especially with square roots, squaring both sides can sometimes lead to other issues, but in this case, since we didn't square the square root term on both sides but isolated it once, it seems fine.  
Let me consider another method to solve the equation, ...



...

Same quadratic as before, so same solutions. So, same process, same result. So,  $x = 1$  and  $x = 5$ , but only  $x = 5$  is valid. But just to make sure, let me consider if squaring the equation introduced any extraneous solutions. Sometimes when we square both sides, it's easy to have extra roots.

...

#### C.4 Collapse of Reasoning Coherence Following Very Long Outputs

Response (DeepSeek-R1-Distill-Qwen-1.5B; budget = 16000 tokens)

...

WaitWait endpoints.  
 about Wait.  
 , say needWait.  
 Maybe Wait Wait set's.  
 me Wait,, mark,:,'s No.  
 the,, me's number minimumButButWait.'s the,.  
 WaitWaitWait.  
 ao over be alone to..  
 , Wait.  
 . the the the Wait Wait,.'s the me's Wait Wait the.'s.  
 me. with of., range's., to (WaitWait k me nonWaitWait me's total  
 number integer's theBut a asagrams, No me,,her.  
 Well Wait Wait a. ...

## D Analyzing Reasoning Traces with LLM-as-a-judge

### D.1 LLM-as-a-judge Setup

In order to illustrate the reasoning structure improvement of SB, we used LLM-as-a-judge to help us analyze the reasoning trace because the lengthy reasoning trace is hard for human to read. We use Gemini 2.5 Flash Preview (Google, 2025) with  $\text{top\_p}=0.95$  and  $\text{temperature}=0.7$ . The prompt for reasoning trace categorization and identifying the first appearance of correct answer as follows:

#### Prompt (Reasoning Categorization)

You are a reasoning trace analyst. Your role is to categorize each line (separated by a newline, but do not treat equations as an individual line) in a given reasoning trace according to the function it serves within the reasoning process. The goal is to understand how different parts of the model's reasoning contribute (or not) to the final answer.

You will be given a **full reasoning trace**, which ends with a final answer after the `</think>` tag. You must:

1. **Only analyze the content before the `</think>` tag**.
2. **Split the reasoning trace into individual lines**, multiple sentences might be in a single line, a line should be ended by a newline.
3. **Assign exactly one label to each line** from the following **mutually exclusive categories**. Choose the **most fitting** label based on the line's function in context.

### Categories:

#### 1. **Adopted Reasoning**

**Definition:** A reasoning step (e.g., setting up a key equation that's shown in the summary, stating a critical definition that's part of the summary, deriving a final value for a variable that's highlighted in the summary, or an intermediate result explicitly carried forward into the summary) whose content or conclusion is **explicitly stated or directly corresponds to a distinct step, equation, or piece of information in the summarized solution presented after the `</think>` tag**. These are the landmark achievements or key building blocks in the reasoning trace that are showcased in the final, condensed explanation.

**Example:** If the final summary after `</think>` includes the equation  $s^2 + 2s - 11.25 = 0$ , then the line in the reasoning trace where this equation is first derived or stated in this form is "Pivotal Reasoning". Similarly, the line stating " $s = 2.5$  km/h" is "Pivotal Reasoning" if  $s = 2.5$  is a key reported value in the final summary.

#### 2. **Exploring Alternatives**

**Definition:** Lines that perform necessary calculations, logical deductions, initial problem interpretation (e.g., "The problem asks for...", "We are given..."), variable definitions (if not themselves pivotal as per the above definition), strategic planning (e.g., "First, I will find  $s$ , then  $t$ ."), or explanations of methodology that **directly contribute to reaching a "Pivotal Reasoning" step or the overall solution, but are not themselves explicitly restated as a distinct element in the final summarized solution**.

**Example:** "Hmm, I need to convert  $t$  into hours because her

walking speed is in km/h." (Problem understanding leading to a pivotal setup). "Multiply both sides by  $s(s+2)$ :  $18 = 1.6 \cdot s(s+2)$ ." (A calculation step leading to a pivotal equation). "Let me compute that left side:  $\dots = 18/(s(s+2))$ " (The process of algebraic simplification).

### 3. **\*\*Repetition & Rephrasing\*\***

\* **\*\*Definition\*\***: Reasoning segments that explore solution paths, hypotheses, methods, or calculations that are **\*\*ultimately abandoned**, identified by the LLM as incorrect or less viable, or otherwise do not form part of the sequence leading to the adopted final solution.**\*\*** This includes trying a different formula that is then discarded, or starting a calculation that is then reversed due to a change in approach.

\* **\*\*Example\*\***: "Perhaps I can solve for  $t$  first by isolating it from equation 1... no, subtracting the equations seems more direct to eliminate  $t$ ." (The statement about solving for  $t$  first by isolation would be "Alternative Path or Exploration" if that path is not taken).

4. **\*\*Verification & Explanation\*\*** \* **\*\*Definition\*\***: Lines that are redundant (e.g., repeating information or conclusions verbatim or near-verbatim shortly after they've been stated, without adding new purpose like emphasis for a next step), purely filler or conversational remarks (e.g., "Okay, so...", "Good.", "Let me think..."), or provide overly trivial explanations/rephrasing that do not advance the solution or significantly aid in understanding the problem-solving process (e.g., "This is a quadratic equation in  $s$ ," immediately after writing  $s^2 + 2s - 11.25 = 0$ ). This also includes restating parts of the problem verbatim without adding interpretation or planning elements, once the initial understanding is established.  
\* **\*\*Example\*\***: "So,  $s$  is 2.5 km/h." (if stated immediately after "Thus,  $s = \dots = 2.5$  km/h" and not serving as a premise for an immediate subsequent calculation). "Let me write them down again:" (if the equations were just presented). "Yes, that seems correct," if it's a low-effort affirmation rather than the concluding part of a rigorous "Verification" step.

### ### Additional Instructions:

\* **\*\*Stop processing as soon as you encounter the '</think>' tag\*\***. Do not categorize anything beyond it. The content after '</think>' is used as the reference for "Pivotal Reasoning". \* If a line **\*\*repeats or paraphrases\*\*** an earlier line without adding new value or serving a clear structural purpose (like summarizing inputs before a new phase of calculation), categorize the repeated instance as **\*\*"Non-Substantive Statement"\*\***, even if the original served a different, valid purpose. \* If a line could arguably fit into more than one category, choose the category that best describes its **\*\*primary function or most specific contribution\*\*** in that context. For instance, a calculation that corrects a previous error is "Verification & Self-Correction" rather than just "Productive Elaboration & Calculation." A statement of a key equation from the final summary is "Pivotal Reasoning" even if it also elaborates. \* Do **\*\*not infer the logical correctness\*\*** of the reasoning or the final answer. The categorization is about the **\*\*structure and an LLM's perceived purpose\*\*** of each statement within its own reasoning

process. \* Treat each line **independently** for categorization, but use the **surrounding context** (preceding and succeeding lines) to understand its function, especially for identifying repetitions, planning statements, or the flow of logic. \* If the reasoning trace starts with the `</think>` tag, i.e.,  
### Output Format:

Return your output as a **JSON array** of objects. Each object should have:

- \* "text": The full original line (string).
- \* "label": One of the five category names above (string).

```
```json [ {"text": "Alternatively, we can try to factor the quadratic equation, but the roots might not be integers.", "label": "Exploring Alternatives" , "text": "Therefore, 
```

Return **only** the structured JSON list, no extra commentary or explanation.

#### Prompt (Identify First Appearance of Correct Answer)

You are a reasoning trace analyst. Your role is to identify the first sentence where the model gets the correct answer of the problem. The goal is to identify the redundant overthinking process after the model has actually solved the problem.

You will be given a full reasoning trace, which ends with a final answer after the `</think>` tag; you will also be given the problem and the correct answer. You must:

1. Only identify the sentence before the `</think>` tag.
2. Identify only the **first** sentence where the model gets the correct answer.
3. The sentence you return should be **exactly the same** as the one in the original reasoning trace.

Return only the sentence you identify, no extra commentary or explanation.

## D.2 Reasoning Category Examples

### Pivotal Reasoning Example

First, I need to figure out the ratio of the enrollments between the two grades. The 7th grade has 520 students, and the 8th grade has 650. So, the total number of students is  $520 + 650$ . Let me calculate that: 520 plus 650 is 1170. So, there are 1170 students in total between the two grades.

### Productive Elaboration & Calculation Example

Therefore, the number of representatives for each grade should be proportional to their enrollment fractions. So, the 7th grade should have  $(4/9)$  of 18 representatives, and the 8th grade should have  $(5/9)$  of 18.

### Exploring Alternatives Example

Hold on, maybe I need to think about the transversal lines and corresponding angles. Since ST is parallel to QR, and PR is a transversal, the angle at T, which is angle STR, should correspond to angle QRP, which is angle R, 105 degrees.

### Verification & Self-Correction Example

Wait, if ST is parallel to QR, then the corresponding angles when PR is the transversal would be equal. So, angle S TR (which is angle STR) should be equal to angle QRP, which is 105 degrees.

### Non-Substantive Statement Example

Wait, maybe I'm confusing the angles. Let me try to think of it differently.



## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [\[Yes\]](#)

Justification: the claims made in abstract/introduction have been elaborated in the rest of the paper (esp. Sections 3 to 5) and in the appendix.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: we have detailed the limitations and future directions in Section 6.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[NA\]](#)

Justification: this paper does not involve formal theorems with assumptions and proof.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

#### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: we have provided the method pipeline in Section 3 and flowchart in Fig. 2. Experiment details are specified in Section 4.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

#### 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: as we mentioned in the beginning of our paper, we will be publicly releasing our codebase. All the datasets (see Section 4) we use are already open-sourced and available online.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: we have provided the most important experimental details in Section 4, and further information in Appendix A.3.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: our experiments involve running an extensive list of 11 benchmarks on 10 LLMs of 1.5B/7B sizes. A single full pass over our testing pipeline (not accounting for the training time) takes over 12 hours on 4 Nvidia A100 GPUs. Therefore, due to the computational constraint, we cannot evaluate and report meaningful statistical significance metrics like error bars in our results.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

#### 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: we have explicitly documented our computational resources in Appendix A.4.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

#### 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: this paper closely follows the NeurIPS Code of Ethics.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

#### 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: we have discussed the broader impact of our paper in the beginning of our appendix.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.

- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: we finetune our models based on the base LLMs described in Section 4, which should have implemented safeguards during pretraining. We manually inspect that our datasets (mainly math and coding benchmarks) do not pose such risks.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: we have properly credited the base models and benchmarks used in Section 4.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.



- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

### 13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: we will provide comprehensive documentations of our data and model when we open-source them.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

### 14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: this paper does not involve crowdsourcing or research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

### 15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: this paper do not involve study participants.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.

- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

#### 16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: in Section 5.2, we have explicitly declared that we have used LLM-as-a-judge as part of our analysis pipeline. We provide further details on the judge LLM, including complete prompts and example outputs, in Appendix D.1.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.